

Course Overview

CS230: System Programming
1st Lecture

Instructors:

Jaehyuk Huh

Overview

- **Course theme**
- **Academic integrity**
- **Five realities**

This course is

- CS230 **System** Programming
- The course is an “introduction to computer systems” from the perspective of programmers
- If you expected System **Programming**, this is not a right course.

Course Perspective

■ Most Systems Courses are Builder-Centric

- Computer Architecture
 - Design pipelined processor
- Operating Systems
 - Implement sample portions of operating system
- Compilers
 - Write compiler for simple language
- Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**

Cheating: Description

■ Please pay close attention

■ What is cheating?

- Sharing code: by copying, retyping, **looking at**, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a lab, line by line
- **Searching the Web for solutions**
- **Using code-generation tools such as ChatGPT**
- Copying code from a previous course or online solution
 - You are only allowed to use code we supply, or from the CS:APP website

■ What is NOT cheating?

- Explaining how to use systems or tools
- Helping others with high-level design issues

■ Ignorance is not an excuse

Cheating: Consequences

■ Penalty for cheating:

- Minimum penalty: one letter grade downgrade (A0 → B0)
- Possible removal from course with failing grade (F)
- Your instructors' personal contempt

■ Detection of cheating:

- We have sophisticated tools for detecting code plagiarism
- In the prior semester, >30 students were caught cheating

■ Commit your updates to your git repository as often as possible

■ Don't do it!

- Start early
- Ask the staff for help when you get stuck

Textbooks

■ Randal E. Bryant and David R. O'Hallaron,

- *Computer Systems: A Programmer's Perspective*, **Third Edition** (CS:APP3e), Pearson, 2016
- <http://csapp.cs.cmu.edu>
- Highly recommend the original English version (you'll have to get used to textbooks in English as soon as possible to survive CS)
- This book really matters for the course! (for labs and exams)
- **Reading the textbook is NOT optional, but required**
 - I will post required sections in KLMS.
 - There will be occasional quizzes about the reading assignments.

■ Optional textbook: Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators

Course Components

■ Lectures

- Higher level concepts

■ Labs (5 assignments)

- **The heart of the course**
- About 2-3 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

■ Exams (midterm + final)

- Test your understanding of concepts & mathematical principles

Getting Help

■ Class Web page: KLMS

- Complete schedule of lectures, exams, and assignments
- Copies of lectures, assignments, exams
- Clarifications to assignments
- **Update your email address (Important announcements can be made via emails and KLMS postings)**

■ Use the **Piazza Q&A** board to share questions and answers

- We will instruct how to use the Piazza board by email
- The Piazza board will be shared by Section A and B.
- Students are also encouraged to answer questions

Getting Help

- **Staff mailing list:** cs230_ta@casys.kaist.ac.kr
 - Use this for all communication with the teaching staff
 - Always CC staff mailing list during email exchanges

Policies: Labs and Exams

■ Work groups

- You must **work alone** on all lab assignments

■ Handins

- Labs due at 11:59pm on Tues or Thurs
- Electronic handins (no exceptions!)

■ Exams

- Midterm + Final (traditional exams)
- **At least 2/3 of exam questions are related to the lab assignments**

■ Appealing grades

- In **writing** to Prof Jaehyuk Huh within 7 days of completion of grading
- Follow formal procedure described in syllabus

Timeliness

■ Lateness penalties

- Get penalized **30% per day**
- No handins later than **1 day after due date**

■ Catastrophic events

- Major illness, death in family, ...
- Formulate a plan to get back on track

■ Advice

- Once you start running late, it's really hard to catch up

Policies: Grading

- **Exams (70%):** midterm (35%), final (35%)
- **Labs (25%):** weighted according to effort
- **Attendance (5%)**
 - Attendance of the first week (8/29 and 8/31) will not be checked as the enrollment changes during the period.
 - If you miss 1/3 of class meetings (9 class meetings), the final grade will be automatically F.

Labs

- ***Tentative*** lab schedule
- Lab 1: Data lab (9/26)
- Lab 2: Bomb lab (10/12)
- Lab 3: Attack lab (11/2)
- Lab 4: Tsh lab (11/16)
- Lab 5: Malloc lab (11/30)

Lab Environments and Requirements

■ Lab system

- A remote linux account will be provided.
- Your submission must be running on the account.

■ Programming environments

- C/C++ is the lab language. If you are not familiar with it, you must learn it in 2 weeks.
- Vscode (Visual Studio code) and git will be used for IDE and repository.
- We will post introductory materials for the environments for self-study.

Programs and Data

■ Topics

- Bits operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

■ Assignments

- L1 (datalab): Manipulating bits
- L2 (bomblab): Defusing a binary bomb
- L3 (attacklab): The basics of code injection attacks

Exceptional Control Flow

■ Topics

- Hardware exceptions, processes, process control, Unix signals
- Includes aspects of compilers, OS, and architecture

■ Assignments

- L4 (tshlab): Writing your own Unix shell.
 - A first introduction to concurrency

Virtual Memory

■ Topics

- Virtual memory, address translation, dynamic storage allocation
- Includes aspects of architecture and OS

■ Assignments

- L5 (malloclab): Writing your own malloc package
 - Get a real feel for systems-level programming

Networking, and Concurrency

■ Topics

- High level and low-level I/O, network programming
- Internet services, Web servers
- concurrency, concurrent server design, threads

Lab Rationale

- **Each lab has a well-defined goal such as solving a puzzle or winning a contest**
- **Doing the lab should result in new skills and concepts**
- **We try to use competition in a fun and healthy way**
 - Set a reasonable threshold for full credit

Course Theme:

Abstraction Is Good But Don't Forget Reality


- **Most CS courses emphasize abstraction**
 - Abstract data types
 - Asymptotic analysis
- **These abstractions have limits**
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- **Useful outcomes from taking cs230**
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!



■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Great Reality #4:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance

- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

*Welcome
and Enjoy!*